

How should I investigate data quality?

Roy Ruddle

Professor of Computing
Director of Research Technology, Leeds Institute for Data Analytics

<https://raruddle.wordpress.com/>



LIDA

UNIVERSITY OF LEEDS

1

Overview

- Introduction
 - A data science workflow
 - Data quality taxonomy
 - What do data scientists want?
- 6-step Data Quality Method
- Software (vizdataquality PyPI package)
- Benefits
 - Save time
 - Reduce re-work
 - Improve results quality



UNIVERSITY OF LEEDS

2

Roy Ruddle



- Inter-disciplinary background
- Expert in data visualization and data quality
 - Hospital and GP data
 - Radio astronomy
 - Supermarkets & convenience stores
 - Adult social care
 - Petrophysics
 - Train doors!

Data quality

- Open source
- Method & PyPI package
- Training




Leeds Virtual Microscope




UNIVERSITY OF LEEDS

3

Workflow



Idealised data science project



Real project



UNIVERSITY OF LEEDS


4

What?

- 69 characterisation and data quality tasks
 - Plain English
 - Each task articulated as question(s) to answer about your data

Task	Question(s)
Example values	What does the data look like (print the first/last few rows)?
Range & variability	What is the minimum/maximum value? What are certain percentiles (e.g., 25th, median and 75th)?
Missing values	How many values are missing? Did you expect the variable to be complete?
Validity	Do any values contain characters or symbols that should not be present?
Violate assumptions?	What assumptions have you made about variables (are they true)?
Frequency measures	What is the distribution of the values? How many times does each value occur?

When: 6-step method



How: vizdataquality

- Python package
 - <https://pypi.org/project/vizdataquality/>
 - 16 example Jupyter notebooks (<https://github.com/royruddle/vizdataquality>)
- Data profiling calculations & visualizations
 - Uses Pandas and Matplotlib
 - If necessary, customise visualizations with **kwargs
- Use the 6-step method to do your investigations
- Clean your data & create a report at the same time
 - Output text, figures & tables as webpage, Latex or for Word

Leeds parking fines

- Open data
 - <https://datamillnorth.org/dataset/v8ggw/off-street-parking-fines>
 - Quarterly datasets, Apr 2013 to Mar 2022
- Example dataset: Jan – Mar 2019
 - See vizdataquality “Workflow (parking fines)” notebook

Step 1: Is anything obviously wrong?

```
df = pd.read_csv('Quarter 4 201819.csv')
```

```
print(df.iloc[-1])
```

PCN LS08541469
 ISSUED 31/03/2019
 LOCATION WEST STREET CP - CITY
 CONTRAVENTION 83 WITHOUT DISPLAYING A VALID TICKET
 FINE 50
 Last Pay Date 31/03/2019
 Total Paid 0.0
 Balance 0.0
 Unnamed: 8 NaN

What time?
 Eh?
 Illogical

Date format is "day first"

UNIVERSITY OF LEEDS


13

Step 2: Watch out for special values

• Watch out!

- Certain values may have a special meaning (e.g., 999999)
- Some systems use '0' as a default value

```
stats = vdqc.calc(df)
cols = ['Balance', 'FINE', 'Total Paid']
fig_kw = {'size_inches': (5, 3.25), 'constrained_layout': True}
ax_kw = {'title': 'Number of zero values', 'xlabel': None,
        'ylim': (0, num_rows), 'ylabel': 'Count'}
vdqp.scalar_bar(stats.loc[cols]['Number of zero values'],
                fig_kw=fig_kw, ax_kw=ax_kw)
```



UNIVERSITY OF LEEDS

14

Step 3: Is any data missing?

• Missing values


- Special values?
- Zeros?
- Invalid values?

• What is the extra "Unnamed: 8" variable?

```
fig_kw = {'size_inches': (5, 3.25), 'constrained_layout': True}
ax_kw = {'title': 'Missing values', 'ylabel': None, 'xlim': (0, num_rows),
        'xlabel': 'Count'}
vdqp.scalar_bar(stats['Number of missing values'], vert=False,
                fig_kw=fig_kw, ax_kw=ax_kw)
```

• Spatial coverage

• Temporal coverage



UNIVERSITY OF LEEDS

15

Step 4: Check each variable

• Uniqueness

• Date & time distributions

• Value lengths


• NB: That's just some of the tasks

UNIVERSITY OF LEEDS

16

Step 4: Uniqueness

```
fig_kw = {'size_inches': (5, 3.25), 'constrained_layout': True}
ax_kw = {'title': 'Unique values', 'xlim': (0, num_rows), 'xlabel': 'Count',
        'ylabel': None}
vdqp.scalar_bar(stats['Number of unique values'], perceptual_threshold=0.02,
               vert=False, datalabels=True, fig_kw=fig_kw, ax_kw=ax_kw)
```



Step 4: So how much was paid in total?

- Some PCNs are paid in parts, not in one go
 - So you will get the wrong answer if you just sum the Total Paid

```
# What is the sum of the Total Paid column?
print('Sum of the Total Paid column =', df['Total Paid'].sum())

# But what was the total amount paid?
print('Total actually paid =', df.groupby('PCN')['Total Paid'].max().sum())


Sum of the Total Paid column = 191106.02
Total actually paid = 106806.93
```

Step 4: Dates

NB: Each PCN is only counted once!

- Check frequency distribution at multiple levels of detail


```
fig_kw = {'size_inches': (12, 3), 'constrained_layout': True}
ax_kw = {'title': 'Distribution of ISSUED values'}
components = ['year', 'month', 'dayofweek']
xrot = [90, 90, 90]
vdqp.plotgrid('datetime distribution', df_noduplicates['ISSUED'], num_rows=1,
              xlabels_rotate=xrot, fig_kw=fig_kw, ax_kw=ax_kw, components=components)
```



Step 4: Value lengths

- How many characters are in the values of each variable?
 - Big ranges often indicate data quality issues

```
fig_kw = {'size_inches': (3, 2)}
ax_kw = {'xlim': 0}
vdqp.dot_whisker(stats['Value lengths'].sort_index(), vert=False,
                 fig_kw=fig_kw, ax_kw=ax_kw)
```



Step 4: Big value length ranges

```
cols = ['LOCATION', 'CONTRAVENTION']
df_examples = vdbc.get_value_lengths_examples(df[cols])

for column in df_examples.columns[1:]:
    print(column)
    for row in df_examples[['Examples', column]].itertuples(index=False):
        print(row[0].rjust(14) + ': ', row[1])
```

LOCATION

Shortest value: 'CASTLE ST CP - CITY'

Median value: 'KIRKSTALL LEISURE CENTRE - CP'

Longest value: 'BARLEY HILL ROAD CP - GARFORTH'

CONTRAVENTION

Shortest value: '73 PARKED WITHOUT PAYMENT'

Median value: '84 ADDITIONAL PAYMENT MADE TO EXTEND TIME'

Longest value: '81 IN A RESTRICTED AREA IN A CAR PARK'



UNIVERSITY OF LEEDS

21

Step 5: Check combinations of variables

- Variety of tasks, e.g.
 - What assumptions have you made?
 - Do any sets of values violate rules?

Variables
PCN
ISSUED
LOCATION
CONTRAVENTION
FINE
Last Pay Date
Total Paid
Balance



UNIVERSITY OF LEEDS

22

Step 5: Violate assumptions

- Why?

Fine - Total Paid - Balance (£)	Count
-116	3
-100	5
-56	4
-50	24
-43	18
-35	28
-33	729
-25	332
-10	1
0	785
25	3352
35	215
49.06	1
49.1	1
50	1559
70	51



UNIVERSITY OF LEEDS

23

Step 6: Profile the cleaned data



UNIVERSITY OF LEEDS

24

Data quality report

- Three formats
 - Text & images (manually import into Word)
 - Latex
 - HTML



UNIVERSITY OF LEEDS

25

Why?

- Do more in less time (know what to do; the tasks & questions)
 - Reduce time & cost (by avoiding re-work)
 - Understand any limitations of the data
 - Correct your assumptions
 - Improve your data
- } Improve results



UNIVERSITY OF LEEDS

26

Summary

- Publicly available resources to help you rigorously investigate data quality and profile your data
 - 6-step Data Quality Method
 - Python package: vizdataquality
- One day CPD course
- YouTube “how to” videos (coming soon)
- Open data repository roll-out



UNIVERSITY OF LEEDS

27

Questions

6-step Data Quality Method

<https://github.com/rovruddle/6-step-data-quality-method>



vizdataquality

<https://pypi.org/project/vizdataquality/>



For these slides, see

<https://github.com/rovruddle/tutorials-and-talks/blob/main/README.md#talks-about-data-quality>



Courses

Birmingham, London, Manchester, Edinburgh:

<https://store.leeds.ac.uk/conferences-and-events/faculty-of-medicine-health/lida-leeds-institute-for-data-analytics>



UNIVERSITY OF LEEDS

28